

Master Thesis Proposal

Algebraic Server Reconciliation for Online Games

Philipp Hinz

Reviewed by Dr.-Ing. Guido Rößling

24. April 2025

1 Introduction

Modern online games often involve multiple players interacting within a shared virtual world. This requires synchronizing game elements, referred to here as „entities,“ across all participants' systems. To combat cheating, a common approach is server-authoritative architecture, where the central server dictates the game's logic and state. Clients transmit their inputs to the server and receive state updates in return. However, network latency, due to physical limitations, introduces a delay in receiving these updates, hindering real-time responsiveness, especially in fast-paced games.

Client-side prediction helps mask this latency. Each client simulates the game's progression based on its inputs, providing immediate feedback. However, this can lead to discrepancies if the client's simulation deviates from the server's authoritative state. Simply overwriting the client's state with the server's is problematic because the server's update reflects a past state. Techniques to reconcile these past state updates with the present client state are known as „server reconciliation.“

2 Related Work

The prevalent server reconciliation method involves rolling back the client's simulation to the past state provided by the server and then reapplying the user's inputs that have occurred since that time. This rollback can be achieved in two primary ways: either the server transmits the complete game state with each update (increasing network load and potential synchronization problems) or, if the game physics are deterministic, the client stores past states and replays other players user inputs.

Both approaches are computationally demanding, as they require the game simulation to run at least twice per frame. Optimizations that avoid full simulation replays often necessitate custom solutions, increasing development complexity and the risk of synchronization errors. Furthermore, the deterministic approach requires each client to possess complete knowledge of the game state, which is not only resource-intensive for large worlds but also increases vulnerability to cheating.

Alternatives to server-authoritative models exist, such as client-authoritative architectures. In this approach, each player has direct control over their character or objects, sending state changes rather than inputs to the server. While easier to implement, this method is highly susceptible to cheating, as players can transmit manipulated state information. Server reconciliation also provides benefit in simulating other player inputs, allowing low-latency feedback.

3 Problem Statement

This thesis introduces **Algebraic Server Reconciliation**, a novel approach to applying past state changes to the current game state. This method aims to reduce computational overhead compared to traditional rollback methods while maintaining flexibility and avoiding the need for deterministic physics. It offers a more general framework than popular existing solutions like GGPO (Good Game Peace Out, a rollback-based networking library).

4 Approach

Algebraic Server Reconciliation leverages the concept of an abelian group to model the game state. This requires defining an associative and commutative addition operation $+$ for the game state, along with a zero element and an inverse (negative) for each state element. The server transmits state **changes** to the client; if no change has occurred, a zero element is sent. The client keeps track of state changes since the last server update.

Upon receiving a server update, the client identifies the corresponding past state change it made at the equivalent time. It then calculates the difference between the server's change and its own past change. This difference represents the divergence between the client's prediction and the server's authoritative state. This difference is then added to the client's current state, effectively correcting the discrepancy. Over time, this process leads to convergence between the client and server states.

More formally:

Let $s_{t_0}^c$ and $s_{t_0}^s$ be the client and server states at time t_0 , respectively. Assume the client has simulated two steps forward, reaching state $s_{t_2}^c = s_{t_0}^c + x_{t_0} + x_{t_1}$, where x_{t_0} and x_{t_1} are the client's state changes at times t_0 and t_1 . The server then sends a state change y_{t_0} . The client updates its state by adding the difference $(y_{t_0} - x_{t_0})$:

$$\begin{aligned} & s_{t_2}^c + (y_{t_0} - x_{t_0}) \\ &= (s_{t_0}^c + x_{t_0} + x_{t_1}) + (y_{t_0} - x_{t_0}) \\ &= (s_{t_0}^c + y_{t_0}) + x_{t_1} + (x_{t_0} - x_{t_0}) \\ &= s_{t_1}^s + x_{t_1} \end{aligned}$$

Assuming $s_{t_0}^c = s_{t_0}^s$

This demonstrates how applying the difference re-integrates the server's past state change without requiring a full rollback and re-simulation.

5 Planned Steps

This thesis will investigate the feasibility and limitations of Algebraic Server Reconciliation, with a particular focus on identifying effective methods for representing game states as abelian groups. The research will explore both the benefits and potential drawbacks of this approach. Two development tracks are planned:

1. **Proof-of-Concept Simulation:** A small-scale simulation will be developed, implementing both Algebraic Server Reconciliation and a conventional rollback-based method. This will enable a direct performance comparison and provide an initial demonstration of the proposed technique's viability.
2. **Full-Fledged Game Prototype:** A more comprehensive game prototype will be constructed to explore the constraints and opportunities presented by a more realistic and complex game environment. This will facilitate the identification of potential limitations and allow for refinement of the approach for practical implementation.

6 Additional Ideas

The core concept of Algebraic Server Reconciliation lends itself to several potential optimizations and refinements:

1. **Entity Removal Handling:** When a client removes an entity (represented as a state change $-a$) before receiving server confirmation, the client must retain a copy of the original entity state (a). This is necessary because the reconciliation process requires calculating $(y - x)$, where x might be $-a$. Therefore, the client needs to be able to compute $-(-a) = a$, effectively „undoing“ the removal locally. Crucially, this full entity data is only required for client-initiated removals. Server-initiated removals, represented by y , only need to contain the entity’s identifier, as the client never needs to compute the inverse of y . This asymmetry helps minimize the data the server needs to transmit.
2. **Efficient Change Comparison with Merkle Trees:** To determine whether client and server state changes for a given entity and its components are identical, Merkle trees can be employed. Instead of comparing the raw data of each change, the client and server can compare the Merkle roots (hashes) of their respective change sets. A mismatch in the roots indicates a discrepancy, while matching roots provide high confidence that the changes are identical, significantly reducing the computational cost of comparison. This is especially beneficial for complex entities with numerous components.

7 (WIP) Proposed solution

The client defines merkel trees over its state changes. Leafs are only needed for nodes where the node content is somewhat larger than the hash (not the case for position). For each sent input we also send the merkel-tree, that is the result of changes caused by this input. One can observe, that we do not lose any time by adding this information. Client side we can compute the merkel-tree in the same moment as the input is captured and server-side we process the input and can instantly compare it to the newly computed merkel-tree. Usually the client wants to send inputs as often as possible, this means, that the server is always very well aware how synced the client is component wise.

The server now usually does not want to send updates each frame but bundeled, especially since we have reconciliation. Intuitively we might thing that we need reliable messages to not miss computations in the algebra, but we dont. The trick is, after sending an update to

8 Modelling

We define a game world using a game state S , an initial game state $s_0 \in S$ and a progression function $f : (S, I) \rightarrow S$ as $G = (S, s_0, f)$ where I is some external input. The game state is of a set of entities $E \subseteq S$. The function f is defined as the folding using constant input I of a list of functions $f_i : I \rightarrow S \rightarrow S$ each individually called a system. A game state at time t can be progressed using some input $i_t \in I$ to time $t + 1$ using the progression function: $s_{t+1} = f(s_t, i_t)$. We can combine the input and state to a frame $r_t = (s_t, i_t)$.

A machine M is running a game G in state $s_t \in S$ at time t . For networking we will look at the following environment, we have the main client $c \in M$, the other client $c' \in M$ and a server (host) $h \in M$. We consider an server authoritative and client / server model. Clients can only communicate with the server and the servers game state is the universal truth. Each client produces some input i_t^m which together gives the input of a frame i_t .

- Game $G = (S, I, s_0, f)$
- Time t
- State $s_t \in S$ at time t
- Input $i_t \in I$ at time t
- Progression function $f : I \rightarrow S \rightarrow S$
- Systems $f_k : I \rightarrow S \rightarrow S$

f is defined as the folding of f_k with constant i_t

$$s_{t+1} = f(s_t, i_t)$$

- Machine $m_t \in M$ running G with state s_t at time t

s_t^m defines the state s_t of machine m_t

- Client $c \in S \subseteq M$ applying input i_t^c

i_t is defined as the collection of all $i_t^c \in i_t$

clients can only communicate with the server

the server holds the universal truth

We have one host h and n clients c_i . Each client c_i has a latency of l_i . Each tick the server sends its state s_t^h to each client, while each client sends their input $i_t^{c_i}$

The server does only progress the state, when all input for one state is received. Each client will receive the initial state with around $\frac{1}{2}l_i$ and is able to respond .

We consider two clients $c_1, c_2 \in C$ and a host $h \in M$. Messages travel with some time, which for simplicity we set to one timestamp. The host starts with a world s_0^h .

In the naive scenario the server continuously sends the state to the clients and the clients send their inputs to the server. The server only progresses the state, when all input arrived. The client will send their first input i_0^c for frame s_0^c . The input will arrive at the server one tick later and the response one tick later again. Therefore the client had to send i_1^c and i_2^c only in response for frame s_0^c before receiving $s_1^c = s_1^h$. As we can see from here on every input i_i^c will be send to the server, needing two ticks before getting a response. Therefore each client will always only be able to respond with an input i_i^c to the state s_{i-2}^c . Since

In the naive scenario the server sends messages the state to the clients, therefore the clients will be at $s_0^c = s_0^h$ while the server is already at s_1^h . The clients respond to the server with their first input arriving at $t = 2$.