

Algebraic Server Reconciliation for Online Games

Algebraische serverseitige Zustandskorrektur für Online-Spiele

Master thesis by Philipp Hinz

Date of submission: October 04, 2042

1. Review: Dr.-Ing. Guido Rößling
Darmstadt



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Erklärung zur Abschlussarbeit gemäß § 22 Abs. 7 APB TU Darmstadt

Hiermit erkläre ich, Philipp Hinz, dass ich die vorliegende Arbeit gemäß § 22 Abs. 7 APB TU Darmstadt selbstständig, ohne Hilfe Dritter und nur mit den angegebenen Quellen und Hilfsmitteln angefertigt habe. Ich habe mit Ausnahme der zitierten Literatur und anderer in der Arbeit genannter Quellen keine fremden Hilfsmittel benutzt. Die von mir bei der Anfertigung dieser wissenschaftlichen Arbeit wörtlich oder inhaltlich benutzte Literatur und alle anderen Quellen habe ich im Text deutlich gekennzeichnet und gesondert aufgeführt. Dies gilt auch für Quellen oder Hilfsmittel aus dem Internet.

Diese Arbeit hat in gleicher oder ähnlicher Form noch keiner Prüfungsbehörde vorgelegen.

Mir ist bekannt, dass im Falle eines Plagiats (§38 Abs.2 APB) ein Täuschungsversuch vorliegt, der dazu führt, dass die Arbeit mit 5,0 bewertet und damit ein Prüfungsversuch verbraucht wird. Abschlussarbeiten dürfen nur einmal wiederholt werden.

Bei einer Thesis des Fachbereichs Architektur entspricht die eingereichte elektronische Fassung dem vorgestellten Modell und den vorgelegten Plänen.

Darmstadt, 04.10.2042

Philipp Hinz



Abstract

This is a template to write your thesis with the corporate design of TU Darmstadt.

Contents

1	Introduction	1
2	Related Work	2
2.1	Foundational Architectures for Latency Mitigation	2
2.1.1	The Authoritative Server and the Challenge of Latency	3
2.1.2	Client-Side Prediction and Server Reconciliation	4
2.1.3	Rollback Netcode	5
2.1.4	Server-Side Lag Compensation	7
2.2	Deterministic Synchronization Models	8
2.2.1	The Deterministic Lockstep Architecture	8
2.2.2	Analysis of Trade-Offs and Application	9
2.3	Theoretical Frameworks for State Convergence	10
2.3.1	Strong Eventual Consistency in Peer-to-Peer Gaming	10
2.3.2	Conflict-Free Replicated Data Types (CRDTs)	10
2.4	Case Studies of Networking in Practice	11
2.4.1	Valve's Source Engine: A Canonical Implementation	12
2.4.2	Overwatch: Determinism in a High-Paced Shooter	12
3	Network model	14
3.1	Overview	14
3.1.1	Naive model	16
3.1.2	Fair naive model	18
3.1.3	Lockstep	20
3.1.4	Player observation limitation	22
3.1.5	Client side prediction	23
3.1.6	Server reconciliation	26
3.1.7	Delta State	26
3.2	Algebraic server reconciliation	27
3.2.1	Definition	27

3.2.2	Abelian group for ECS games	28
3.2.3	Limitations	28
4	Networking Architecture	29
4.1	Overview	30
4.1.1	Four quadrants model	31
4.1.2	Server and Client	32
4.1.3	Game and Network	33
4.2	Server networking layer	35
4.3	Client networking layer	36
4.3.1	Rollback reconciliation	40
4.3.2	Algebraic reconciliation	42
5	Experiments	44
6	Discussion	45
7	Future work	46



1 Introduction

2 Related Work

The challenge of maintaining a consistent and responsive shared reality across a distributed system is a foundational problem in computer science. In the domain of real-time networked multiplayer games, this challenge is amplified by stringent latency requirements and the demand for interactive fairness. This chapter surveys the landscape of techniques developed to address state synchronization in this context. A notable characteristic of this field is the bifurcation of knowledge dissemination; while foundational principles of distributed computing are well-documented in academic literature, many state-of-the-art implementations and architectural patterns are primarily discussed within industry forums, such as the Game Developers Conference (GDC), and documented in non-peer-reviewed materials like engine wikis and developer blogs. This is a direct consequence of the commercial, fast-paced, and often proprietary nature of game development, where competitive advantage can hinge on networking solutions. Consequently, this review adopts a synthetic approach, integrating formal academic research with seminal industry contributions to provide a holistic overview. The chapter begins by examining the foundational architectures for latency mitigation in the prevalent client-server model. It then contrasts this with the alternative paradigm of deterministic synchronization. Subsequently, it explores theoretical frameworks from distributed systems research, such as Conflict-Free Replicated Data Types (CRDTs), that offer alternative models for state convergence. Finally, it grounds these discussions in case studies of influential, real-world implementations.

2.1 Foundational Architectures for Latency Mitigation

To combat the unresponsiveness caused by network latency, modern games don't eliminate delay but instead conceal it. This chapter introduces the foundational techniques that create this illusion of instantaneous interaction. We will cover the core

strategies: Client-Side Prediction for immediate feedback, Server Reconciliation to correct errors, and Server-Side Lag Compensation to ensure fairness.

2.1.1 The Authoritative Server and the Challenge of Latency

The predominant architecture for modern multiplayer games is the client-server model, wherein a single server is designated as the authoritative source of truth for the game state. Clients send user inputs to the server and, in return, receive periodic updates, or “snapshots,” of the world state. This model centralizes simulation and prevents many forms of cheating, as the client is never fully trusted with game-critical logic. However, it introduces the fundamental problem of network latency: the round-trip time (RTT), illustrated in Figure 1, which is the total time it takes for a client’s input to reach the server and for the server’s response to return. For fast-paced action games, where delays of even a few milliseconds can be perceptible, this latency renders a naive implementation unplayable, as there would be a significant delay between a player’s action and its visual feedback. The following sections detail the primary techniques developed to conceal or compensate for this inherent latency.




Figure 2.1: Illustration of network latency between a host and a client. The round-trip time (RTT) is the sum of the send and receive latencies.

2.1.2 Client-Side Prediction and Server Reconciliation

To combat the feeling of unresponsiveness, the most widely adopted technique is Client-Side Prediction. The client does not wait for the server's confirmation to execute a local player's command. Instead, it speculatively executes the input and simulates the outcome locally, providing immediate visual feedback. This creates the illusion of a zero-latency environment for the local player's movement and actions. The earliest known implementation in a first-person shooter was in Duke Nukem 3D (1996), as shown in Figure 2, with the technique being popularized by id Software's QuakeWorld.



Figure 2.2: A screenshot from Duke Nukem 3D (1996). The game is one of the earliest known first-person shooters to implement client-side prediction, providing players with immediate visual feedback for their movement to hide network latency.

This speculative execution, however, introduces the problem of divergence. The client's predicted state may differ from the server's authoritative state, which has perfect information about the world, including the actions of other players and collision physics. This discrepancy is known as a prediction error. To resolve this, a process called Server Reconciliation is employed. When the client receives an authoritative state update from the server, it compares this state with its own predicted state. If a mismatch is detected, the client must correct its state to match the server's. A naive correction would result in a jarring visual “snap” as the player character is teleported to the correct position. To mitigate this, a crucial step is performed: the client replays all the inputs that were sent

to the server but have not yet been acknowledged in the received server snapshot. By reapplying these inputs from the corrected server state, the client can re-simulate its way back to the present, often resulting in a correct final position and eliminating the visual artifact entirely in most cases. This process requires clients to send inputs with a sequence number and for the server to echo back the number of the last processed input.

2.1.3 Rollback Netcode

Rollback netcode can be understood as a more aggressive and generalized application of the principles of prediction and reconciliation, primarily developed to meet the stringent timing requirements of fighting games. While traditional client-side prediction is typically applied only to the local player, rollback netcode predicts the inputs of the remote player as well, often by assuming they will repeat their last known input. The game simulation for all players advances based on these local and predicted inputs. When an actual input from the remote player arrives, the game state is ‘rolled back’ to the frame before the prediction began, the correct input is inserted, and the simulation is ‘fast-forwarded’ back to the present time. This multi-step correction process is visualized in Figure 3. The entire rollback and re-simulation ideally occurs within a single frame, making the correction imperceptible to the user.




Figure 2.3: Illustration of the rollback netcode process. (a) The client's initial predicted state progression. (b) The server's authoritative state diverges (State G instead of B). (c) Upon receiving the authoritative State G, the client rolls back to State A, applies State G, and then re-simulates its subsequent inputs to reach States H and I. (d) The server again processes inputs and generates a new authoritative State B, which differs from the client's expectation (G). (e) The client performs another rollback, reverting to a prior known good state, applying the server's authoritative State B, and re-simulating all inputs to reach State C. This demonstrates how rollback continuously adjusts the client's timeline by repeatedly re-computing future states (e.g., the state at time point 4 is computed multiple times) based on new authoritative information, ensuring eventual consistency.

The primary benefit is a highly responsive experience that feels akin to offline play, even under significant latency. However, the implementation is complex, requiring the ability to save and load the complete game state very rapidly for every frame and to re-simulate multiple frames' worth of logic in a fraction of a second. While rollbacks aim to be imperceptible, extreme cases or visual debuggers can reveal the momentary “snapping” as shown in Figure 4. Michael Stallone of NetherRealm Studios noted that retrofitting rollback into Mortal Kombat X took approximately two man-years of effort, with significant challenges in optimizing serialization and ensuring non-deterministic elements like visual and audio effects were handled consistently during rollbacks.




Figure 2.4: A screenshot from the fighting game Mortal Kombat X (2015). The genre’s demand for frame-perfect inputs makes it extremely sensitive to latency, establishing fighting games as a primary driver for the adoption of rollback netcode. The challenges of retrofitting this technology into Mortal Kombat X are a well-documented example of its implementation complexity.

2.1.4 Server-Side Lag Compensation

While client-side prediction addresses the local player’s sense of responsiveness, it does not solve the problem of interactional fairness. A player might fire at a target that is clearly visible on their screen, but due to latency, the target has already moved to a new position

on the server's timeline by the time the shot command is processed. To solve this, servers employ Lag Compensation. The server maintains a short history of past player positions. When it receives a user command, such as a shot, it uses the packet's latency to estimate the time at which the command was actually executed on the client. It then temporarily "rewinds" the positions of other players in the world to where they were at that moment in the past. The hit detection is then performed against these historical positions. After the command is processed, the players are returned to their current positions.

This technique, famously implemented in the Source Engine, is often described as "favor the shooter". It ensures that if a player sees a valid target, their shot will register, creating a more reliable and fair experience for the attacker. However, it can lead to perceptual paradoxes for the victim, who may be hit after they have already taken cover on their own screen. This is an accepted trade-off in most modern shooters, prioritizing interactional consistency over strict temporal accuracy. The collection of these techniques reveals a core design duality in networked games. One branch of techniques, encompassing prediction, reconciliation, and rollback, is focused on optimizing the local player's subjective feeling of control and responsiveness. The other branch, exemplified by lag compensation, is focused on optimizing the acting player's objective sense of fairness during interactions. These two goals can be in direct conflict, and the specific balance chosen by a game's networking model defines much of its "feel."

2.2 Deterministic Synchronization Models

In sharp contrast to sending frequent state updates, deterministic models synchronize only player inputs. This section explores architectures where every client runs an identical simulation. Provided each client processes the same inputs in the same order, their worlds remain perfectly synchronized. We will examine the classic lockstep model, its profound bandwidth advantages, and its demanding trade-offs, such as the strict requirement for perfect determinism.

2.2.1 The Deterministic Lockstep Architecture

In contrast to the state-synchronization model employed by authoritative servers, the deterministic lockstep model operates on a fundamentally different principle. Instead of transmitting game state, clients only transmit their inputs to one another, either through a peer-to-peer topology or relayed via a simple server. Each client runs an identical, fully

deterministic simulation of the game world. Provided that every client starts from the same initial state and processes the exact same sequence of inputs in the exact same order (at the same “tick”), their game states are guaranteed to remain perfectly synchronized.

To ensure inputs are processed in the same order, the simulation proceeds in discrete turns. The game will not advance to the next turn until it has received the inputs from all players for that turn. To hide the latency of waiting for these inputs, games implementing this model introduce a fixed input delay, buffering inputs for a short period before executing them. This model’s primary advantage is its extraordinary bandwidth efficiency concerning the number of game entities. Since only small input packets are transmitted, the network load scales with the number of players, not the number of units on screen, making it ideal for genres with potentially thousands of entities, such as Real-Time Strategy (RTS) games.

2.2.2 Analysis of Trade-Offs and Application

The benefits of the lockstep model come with severe trade-offs. The entire system is only as fast as the slowest participant; if one player’s input packet is delayed, the simulation for all players must freeze until it arrives. This makes the model highly susceptible to players with high latency or unstable connections and limits its practical application to games with a relatively low player count.

Furthermore, the requirement of perfect determinism is a formidable software engineering challenge. The simulation must produce bit-identical results across different hardware, operating systems, and compiler settings. This precludes the use of common non-deterministic elements, most notably standard floating-point arithmetic, which can yield slightly different results on different processor architectures. Developers must rely on fixed-point math and carefully control sources of randomness and data structure ordering to prevent desynchronization, or “desyncs,” which are catastrophic failures in this model. A modern evolution, deterministic rollback, combines this model with the rollback techniques described earlier, predicting inputs to avoid freezing and resimulating when the actual inputs arrive.

The choice between an authoritative server model and a deterministic lockstep model represents a fundamental architectural decision to shift the primary burden of complexity. The authoritative server model places complexity in the network layer: managing high bandwidth, compressing state, interpolating between snapshots, and reconciling prediction errors. The deterministic lockstep model places complexity in the application layer: enforcing perfect determinism, managing fixed-point math, and dealing with the

game design implications of input delay and simulation freezes. The state-sync model effectively says, “Let the application be complex and non-deterministic; we will solve the resulting inconsistencies at the network level with authority and correction.” The input-sync model says, “Let’s make the network problem trivial by enforcing extreme constraints on the application itself.” The choice depends entirely on the game’s genre and design priorities.

2.3 Theoretical Frameworks for State Convergence

Beyond bespoke gaming solutions, the field of distributed systems provides formal frameworks for state synchronization. This section introduces these theoretical concepts, focusing on Strong Eventual Consistency (SEC) and a powerful data structure that enables it: the Conflict-Free Replicated Data Type (CRDT). We will explore how CRDTs mathematically guarantee that disparate clients will eventually reach the same state, offering a robust model for decentralized or peer-to-peer architectures.

2.3.1 Strong Eventual Consistency in Peer-to-Peer Gaming

Beyond the specific implementations in gaming, the field of distributed systems provides a theoretical foundation for reasoning about state convergence. A key concept is Strong Eventual Consistency (SEC), which guarantees that if a set of replicas has received the same set of updates, they will be in the same state. Unlike strong consistency models which may require coordination or locking before an operation can be performed, eventual consistency allows replicas to be updated independently and asynchronously, with convergence guaranteed over time. This model is highly relevant for decentralized or peer-to-peer (P2P) game architectures, which lack a central authoritative server to resolve conflicts. The challenge lies in designing data structures and operations that can guarantee this convergence property automatically.

2.3.2 Conflict-Free Replicated Data Types (CRDTs)

Conflict-Free Replicated Data Types (CRDTs) are a class of data structures designed to provide SEC. They allow for concurrent updates on different replicas without coordination, and provide a mathematically proven guarantee that the replicas will

converge. This is achieved by designing operations that possess specific algebraic properties.

There are two primary forms of CRDTs: State-based CRDTs (Convergent Replicated Data Types, or CvRDTs): Each replica can send its entire state to another replica. The receiving replica merges the incoming state with its own using a merge function. To guarantee convergence, this merge function must be associative, commutative, and idempotent. A set of states and a merge function with these properties form a mathematical structure known as a join-semilattice.

Operation-based CRDTs (Commutative Replicated Data Types, or CmRDTs): Instead of shipping the full state, replicas broadcast the update operations themselves. To ensure convergence, these operations must be commutative, meaning they can be applied in any order and still yield the same result. This approach is more bandwidth-efficient but often requires more stringent guarantees from the network layer, such as ensuring that operations are not dropped or duplicated.

While CRDTs are successfully used in collaborative editing software and distributed databases, their application to complex game states is an active area of research. A primary challenge is that many game-world interactions are inherently non-commutative; for example, a Move operation followed by a Shoot operation has a different outcome than the reverse. This creates a significant “impedance mismatch” between the clean, algebraic properties required by simple CRDTs and the messy, stateful, and order-dependent nature of game logic. Applying CRDTs to gaming therefore requires either carefully designing game mechanics to be commutative or developing more complex data structures and algorithms that can handle these ordering constraints, which can erode the initial simplicity of the model. The complexity is not eliminated; it is moved from the network protocol into the data modeling of the game itself.

2.4 Case Studies of Networking in Practice

To bridge the gap between theory and implementation, this section examines the networking models of influential, real-world games. By analyzing seminal titles, we can see how the previously discussed concepts are applied to solve practical challenges. We will dissect Valve’s Source Engine as a canonical example of client-side prediction and lag compensation, and Blizzard’s Overwatch as a modern hybrid that leverages determinism within an authoritative server to achieve high-fidelity competitive play.

2.4.1 Valve's Source Engine: A Canonical Implementation

The networking model of Valve's Source Engine, which powers titles like Half-Life 2 and Counter-Strike: Source, serves as a well-documented and highly influential example of the authoritative client-server architecture. Its documentation, made widely available to the modding and development community, has codified many of the foundational techniques discussed in this chapter. The engine simulates the game in discrete “ticks” on an authoritative server. Clients run prediction for the local player's movement to mask latency, and the server performs lag compensation to ensure fairness in hit registration by rewinding player history. The client also performs interpolation between received server snapshots to produce smooth motion for remote entities, deliberately rendering the world slightly in the past to ensure it always has valid states to interpolate between. The extensive set of user-configurable variables (`cl_interp`, `cl_cmdrate`, `rate`) allows for fine-tuning of the trade-offs between responsiveness and smoothness, exposing the underlying mechanics to the end-user. The Source Engine's networking model represents a robust, battle-tested blueprint for state synchronization in fast-paced first-person shooters.

2.4.2 Overwatch: Determinism in a High-Paced Shooter

Blizzard Entertainment's Overwatch presents a more modern and novel architecture, as detailed in a GDC 2017 presentation by Timothy Ford. While it operates on an authoritative client-server model, it uniquely “leverages determinism to achieve responsiveness and precision”. This represents a hybrid approach that synthesizes principles from both the state-sync and input-sync paradigms. The server runs the simulation at a fixed tick rate (typically 60Hz for competitive play) and is the final authority on game state. However, the simulation logic itself is deterministic. The client is able to run the exact same simulation code as the server.

This design has significant benefits for client-side prediction. Because the client's simulation is identical to the server's, prediction errors are not caused by subtle divergences in physics or logic (e.g., floating-point inaccuracies). Instead, prediction errors arise only from a lack of information—namely, the inputs of other players that have not yet been received. When the server sends a correction, the client can reconcile its state with high confidence, knowing that the underlying simulation logic is sound. This hybrid model aims to achieve the predictive accuracy of a deterministic system within the responsive, cheat-resistant framework of an authoritative server. This signifies a convergence of the two historically separate paradigms, demonstrating that determinism



is not an all-or-nothing architectural choice but can be employed as a powerful tool within a traditional client-server framework to dramatically improve the quality of prediction and reconciliation. The game's netcode also features an “adaptive interpolation delay” system, which attempts to dynamically adjust the rendering buffer based on network conditions to keep gameplay smooth.

3 Network model

To rigorously discuss and analyze networking techniques for online games, particularly those involving state synchronization and reconciliation, it is essential to first establish a clear and formal understanding of the underlying concepts. This chapter lays the theoretical groundwork by defining fundamental components of a game from a networking perspective, such as game state, state progression, and the roles of clients and the server. Building upon these definitions, we will then explore several established network models, from naive approaches to more sophisticated techniques like client-side prediction and server reconciliation. This exploration will highlight the challenges inherent in achieving responsive and consistent gameplay in a networked environment. Finally, this chapter will introduce the core theoretical concepts behind delta states and culminate in the formal definition of algebraic server reconciliation, the novel method investigated in this thesis.

3.1 Overview

We define a game world using a game state S , an initial game state $s_0 \in S$ and a progression function $f : (S, I) \rightarrow S$ as $G = (S, s_0, f)$ where I is some external input. A game state at time t can be progressed using some input $i_t \in I$ to time $t + 1$ using the progression function: $s_{t+1} = f(s_t, i_t)$. We can combine the input and state to a frame $r_t = (s_t, i_t)$.

We say that a machine $m \in M$ is running the game G with state s_t at some time t . Clients are machines $c \in C \subseteq M$ which contribute input $i_t^c \in I^c$ to form the whole input $i_t \in I$. We assume that one machine is designated as the host and all clients can only communicate with the host. The state on the host is considered the truth if states between machines differ.

Example of a formal game definition

To illustrate the concepts introduced thus far, we consider a simplified game inspired by the mechanics of “agar.io”. In this game, the world is viewed from a top-down perspective where each player is represented as a circle. While the original game features collision and consumption mechanics, our simplified version focuses solely on movement to clearly demonstrate the networking concepts. Each player is assigned a unique identifier to distinguish them within the game state.

We formally define this “top-down movement game” as follows:

Vector	$V \subseteq \mathbb{R} \times \mathbb{R}$
Player	$P = V$
State	$S \subseteq \text{id} \times P$
Input for a single player	$I_P \subseteq \mathbb{B}^4$
Input	$I \subseteq \text{id} \times I_P$

The input representation captures the four directional keys that can be pressed during a single frame. Assuming standard addition operations on boolean values and vectors, we define the progression function:

$$\begin{aligned} f^{\text{transform}}((i_{\text{up}}, i_{\text{down}}, i_{\text{left}}, i_{\text{right}})) &= ((i_{\text{down}} - i_{\text{up}}, i_{\text{right}} - i_{\text{left}})) \\ f^{\text{find}}(\text{id}, S) &= x_P, \quad x_P \in \{(\text{id}', x') \in S \mid \text{id} = \text{id}'\} \\ f(s, i) &= \{(\text{id}, s_p + f^{\text{transform}}(f^{\text{find}}(\text{id}, i))) \mid (\text{id}, s_p) \in s\} \end{aligned}$$

Each player consists of a two-dimensional position vector. The game state comprises the set of all player positions indexed by their identifiers. The progression function operates in two stages: first, it extracts the appropriate input for each player from the combined input set, then it transforms these boolean directional inputs into velocity vectors that update the player positions.

Consider the following concrete example with two players:

$$\begin{aligned} s_0 &= \{(p_0, (0, 0)), (p_1, (5, 5))\} \\ i_0 &= \{(p_0, (1, 0, 0, 1)), (p_1, (0, 0, 1, 0))\} \\ f(s_0, i_0) &= s_1 = \{(p_0, (0 + 1, 0 + 1)), (p_1, (5 - 1, 5))\} \\ &= \{(p_0, (1, 1)), (p_1, (4, 5))\} \end{aligned}$$

In this example, player p_0 moves diagonally (up and right) while player p_1 moves left, demonstrating how discrete boolean inputs translate into continuous position updates.

Example of client-host architecture

Consider a multiplayer session with two clients c_0 and c_1 connected to a host h . Using our previously defined movement game, we examine how inputs from multiple clients are processed and synchronized through the central host.

Client c_0 generates input $i_t^{c_0} = (1, 0, 0, 1)$, indicating upward and rightward movement. Simultaneously, client c_1 produces input $i_t^{c_1} = (0, 0, 1, 0)$, representing leftward movement. For multiplayer functionality to work correctly, each client must be aware of all other players' inputs, necessitating a mechanism for input distribution.

Our architecture employs a central host model for this distribution. Following the principle that client-side computation cannot be trusted for security reasons, only the host possesses the authority to evaluate the progression function and determine the canonical game state. This does not preclude clients from performing local predictions for responsiveness, as we shall explore in later sections.

The communication flow proceeds as follows: both clients continuously transmit their inputs to the host h . Upon receiving these inputs, the host updates its authoritative state using the progression function:

$$f(s_0, \{(p_0, i_t^{c_0}), (p_1, i_t^{c_1})\}) = s_1$$

The host then distributes the updated state to all clients. While various optimization strategies exist for this distribution (such as sending only state deltas or relevant subsets), the simplest implementation transmits the complete state s_1 to ensure all clients maintain a consistent view of the game world.

3.1.1 Naive model

We have one host h and n clients c_i . Each client c has a round trip time of t_c^r . We define the round trip time as the sum of the send t_c^s and receive t_c^e latency $t_c^r = t_c^s + t_c^e$. At each tick, the host sends its state s_t to each client and each client sends input i_t^c .

Both the client and the host run with a fixed time interval t^i which is usually smaller than t_c^r . Therefore the host cannot wait with processing all input i_t^c and will use the most recent

input instead. This means that each client will only be able to react with input i_t^c to state $s_{t-t_c^r}$. We call this the naive model.

Example of naive model networking

To demonstrate the temporal challenges inherent in the naive model, we trace through a concrete execution scenario using our movement game. Consider two clients with asymmetric network conditions:

- Client c_0 : round trip time $t_{c_0}^r = 50$ ms (25 ms send, 25 ms receive)
- Client c_1 : round trip time $t_{c_1}^r = 200$ ms (100 ms send, 100 ms receive)
- Fixed tick interval: $t^i = 50$ ms

Starting with initial state $s_0 = \{(p_0, (0, 0)), (p_1, (5, 5))\}$, the host broadcasts this state at tick 0. The following table traces the first 10 ticks of execution:

Tick	Host State	c_0 Sees	c_1 Sees	c_0 Input	c_1 Input	Inputs at Host
0	(0, 0), (5, 5)	—	—	—	—	None
1	(0, 0), (5, 5)	—	—	—	—	None
2	(0, 0), (5, 5)	(0, 0), (5, 5)	—	right	—	None
3	(0, 1), (5, 5)	(0, 0), (5, 5)	(0, 0), (5, 5)	right	—	c_0 : right
4	(0, 2), (5, 5)	(0, 1), (5, 5)	(0, 0), (5, 5)	right	up	c_0 : right
5	(0, 3), (5, 5)	(0, 2), (5, 5)	(0, 1), (5, 5)	right	up	c_0 : right
6	(0, 4), (6, 5)	(0, 3), (5, 5)	(0, 2), (5, 5)	right	up	Both inputs
7	(0, 5), (7, 5)	(0, 4), (6, 5)	(0, 3), (5, 5)	right	up	Both inputs
8	(0, 6), (8, 5)	(0, 5), (7, 5)	(0, 4), (6, 5)	right	up	Both inputs
9	(0, 7), (9, 5)	(0, 6), (8, 5)	(0, 5), (7, 5)	right	up	Both inputs

The execution reveals two critical issues with the naive model. First, examining the input delay: client c_0 receives the initial state at tick 2 and immediately sends a “right” input. This input reaches the host at tick 3, where it gets processed and updates player p_0 to position (0, 1). However, c_0 doesn’t observe this change until tick 4, a full round trip time of 50 ms after sending the input.

Similarly, client c_1 receives the initial state at tick 3, sends an “up” input at tick 4, but the host doesn’t process it until tick 6 due to the 100 ms transmission delay. Client c_1

finally sees their first movement at tick 8, experiencing a 200 ms delay between action and feedback.

The second, more subtle issue is the competitive unfairness in control responsiveness. Client c_0 observes their input taking effect at tick 4, allowing them to make informed decisions about their next move by tick 5. If they realize they're moving in the wrong direction, they can correct course with only a 50 ms feedback loop. In contrast, client c_1 doesn't see their first movement until tick 8, a 200 ms delay that severely hampers their ability to make tactical adjustments.

Consider a scenario where both players accidentally move toward a hazard: client c_0 can recognize and correct their mistake within 2 ticks, while client c_1 continues blind movement for 4 ticks before seeing any feedback. This creates a fundamental inequality where the lower-latency player has superior control precision, faster error correction, and more responsive gameplay, advantages that compound over time and cannot be overcome through player skill alone.

3.1.2 Fair naive model

Each client c can only respond to a state s_t with an input arriving at the host t_c^r time later. Since t_c^r is different for every player, this is not fair.

We define a fair model as: given a fixed delay δ , every client can respond to the same state s_t with input $i_{t+\delta}^c$. We can make the previous model fair by only processing input on the host, once we received input from all players. Every player is able to react to state s_t with the same delay of $\delta = \max_i t_i^r$. Intuitively, we artificially delay what each player can see to the slowest player. We call this the fair naive model.

Many games are based on reaction, doing an action in response to an event. A major problem with the (fair) naive model is that player inputs in reaction to a present state will only be applied to a future state on the host.

Example of fair naive model networking

The fair naive model addresses the competitive inequality of the naive model by synchronizing all players to the slowest connection. Using the same network configuration and initial state:

- Client c_0 : round trip time $t_{c_0}^r = 50$ ms (25 ms send, 25 ms receive)

- Client c_1 : round trip time $t_{c_1}^r = 200$ ms (100 ms send, 100 ms receive)
- Fixed tick interval: $t^i = 50$ ms

Starting with initial state $s_0 = \{(p_0, (0, 0)), (p_1, (5, 5))\}$, the host broadcasts at tick 0. The key difference is that the host now buffers inputs until all players' commands for a given game state have arrived:

Tick	Host State	c_0 Sees	c_1 Sees	c_0 Input	c_1 Input	Host Buffers
0	(0, 0), (5, 5)	—	—	—	—	Empty
1	(0, 0), (5, 5)	—	—	—	—	Empty
2	(0, 0), (5, 5)	(0, 0), (5, 5)	—	right	—	Empty
3	(0, 0), (5, 5)	(0, 0), (5, 5)	(0, 0), (5, 5)	right	—	c_0 : right
4	(0, 0), (5, 5)	(0, 0), (5, 5)	(0, 0), (5, 5)	right	up	c_0 : right
5	(0, 0), (5, 5)	(0, 0), (5, 5)	(0, 0), (5, 5)	right	up	c_0 : right
6	(0, 1), (6, 5)	(0, 0), (5, 5)	(0, 0), (5, 5)	right	up	Both active
7	(0, 2), (7, 5)	(0, 0), (5, 5)	(0, 0), (5, 5)	right	up	Both active
8	(0, 3), (8, 5)	(0, 1), (6, 5)	(0, 0), (5, 5)	right	up	Both active
9	(0, 4), (9, 5)	(0, 2), (7, 5)	(0, 1), (6, 5)	right	up	Both active

The fair naive model fundamentally changes how the host processes inputs. When client c_0 's input arrives at tick 3, the host does not immediately apply it. Instead, it stores this input in a buffer and continues waiting. The host maintains the game state at (0, 0), (5, 5) for ticks 3, 4, and 5, despite having valid input from c_0 . Only at tick 6, when client c_1 's input finally arrives after its 100 ms journey, does the host process both inputs simultaneously.

This buffering strategy creates perfect fairness in reaction timing. Both clients receive the initial state and have the opportunity to respond before any inputs are processed. When the host finally applies inputs at tick 6, it creates state (0, 1), (6, 5) where both players have moved. Client c_0 sees this result at tick 8, while client c_1 sees it at tick 9. Critically, both players observe their inputs taking effect relative to the same initial state, eliminating the advantage that lower latency previously provided.

However, this fairness comes at a significant cost to responsiveness. Client c_0 , who previously enjoyed a 50 ms feedback loop, now experiences a 200 ms delay between input and observation, matching the slowest player. The buffering essentially forces all players

to experience the worst-case latency in the session. In our example, client c_0 sends their input at tick 2 but doesn't see the result until tick 8, a delay of 300 ms total.

This artificial throttling makes the game feel sluggish for players with good connections, potentially driving them away from servers with high-latency participants. The fair naive model thus trades individual responsiveness for competitive balance, a compromise that works poorly for fast-paced action games where immediate feedback is crucial to the gameplay experience.

3.1.3 Lockstep

In the fair naive model, the host is applying old inputs because the interval time is fixed and usually faster than t_c^r . We can loosen this requirement and assume a dynamic interval time. Players only send input i_t^c , after receiving s_t . The host only progresses to state s_{t+1} when all i_t^c have arrived. We now have a fair game while giving players the most recent information to react to.

This model is known as the lockstep model and popular in strategy games, where latency or a longer interval are not a big problem.

Example of lockstep model

The lockstep model ensures perfect synchronization by pausing the entire simulation after each state update until all clients have received and responded. Using the same network configuration:

- Client c_0 : round trip time $t_{c_0}^r = 50$ ms (25 ms send, 25 ms receive)
- Client c_1 : round trip time $t_{c_1}^r = 200$ ms (100 ms send, 100 ms receive)
- Fixed tick interval: $t^i = 50$ ms

Starting with initial state $s_0 = \{(p_0, (0, 0)), (p_1, (5, 5))\}$, the host initiates the first lockstep cycle at tick 0:

Tick	Host State	c_0 Sees	c_1 Sees	c_0 Input	c_1 Input	Host Action
0	$(0, 0), (5, 5)$	—	—	—	—	Broadcast s_0
1	Waiting	$(0, 0), (5, 5)$	—	right	—	Wait for all
2	Waiting	Sent right	$(0, 0), (5, 5)$	—	up	Wait for all

3	Waiting	Waiting	Sent up	—	—	Wait for all
4	(0, 1), (6, 5)	Waiting	Waiting	—	—	Process inputs
5	(0, 1), (6, 5)	(0, 1), (6, 5)	Waiting	right	—	Broadcast s_1
6	Waiting	Sent right	(0, 1), (6, 5)	—	up	Wait for all
7	Waiting	Waiting	Sent up	—	—	Wait for all
8	(0, 2), (7, 5)	Waiting	Waiting	—	—	Process inputs
9	(0, 2), (7, 5)	(0, 2), (7, 5)	Waiting	right	—	Broadcast s_2

The lockstep model operates in discrete, synchronized cycles that guarantee perfect consistency across all clients. Each cycle follows a rigid pattern: the host broadcasts a state, waits for all clients to receive it, waits for all clients to send their inputs based on that state, and only then processes those inputs to generate the next state. This creates natural synchronization points where the entire game pauses until the slowest participant catches up.

In our execution, the first cycle begins at tick 0 when the host broadcasts s_0 . Client c_0 receives this state at tick 1 and immediately sends a “right” input that reaches the host at tick 2. However, the host cannot proceed because client c_1 only receives the state at tick 2 and sends an “up” input that doesn’t arrive until tick 4. Only then can the host process both inputs together, updating the state to (0, 1), (6, 5) and beginning the second cycle at tick 5.

The lockstep model achieves perfect fairness and eliminates all prediction or synchronization issues. Every client sees exactly the same game state before making decisions, and all inputs are processed simultaneously with no player gaining an advantage from lower latency. This makes it ideal for turn-based strategy games or scenarios where absolute consistency matters more than responsiveness.

However, the cost is severe: the game effectively runs at the speed of the slowest connection. With client c_1 ’s 200 ms round trip time, the game can only update every 4 ticks (200 ms), reducing the effective update rate from 20 Hz to 5 Hz. This dramatic slowdown makes the lockstep model unsuitable for any real-time game. Furthermore, if any client experiences a temporary network spike or packet loss, the entire game freezes for all players until that client recovers. The model essentially chains all players to the worst network conditions present in the session, making it impractical for modern action games where fluid, responsive gameplay is essential.

3.1.4 Player observation limitation

We are bound in the time a player is able to react to other players' inputs. One player has to first send their inputs and another player can only then receive them. Any changes to the state from player c' will never be visible to another player c faster than $t_c^e + t_{c'}^s$.

When the interval rate is fixed to a time $\delta < t_c^e + t_{c'}^s$, which is usually the case, a client c will already have made an input i_{t+1}^c without there being a possibility of receiving input $i_t^{c'}$ from client c' . Therefore according to our definition, games with a fixed interval cannot be fair.

Example of player observation limitation

To illustrate this fundamental constraint, consider our two clients attempting to react to each other's movements. Assume optimal conditions with immediate processing and no artificial delays:

- Client c_0 : send 25 ms, receive 25 ms
- Client c_1 : send 100 ms, receive 100 ms
- Fixed tick interval: $t^i = 50$ ms

Suppose at tick 0, player p_1 suddenly moves to block a pathway at position (3, 5). Client c_1 sends this input immediately, taking 100 ms to reach the host. The host processes it instantly and broadcasts the update, which takes another 25 ms to reach client c_0 . Therefore, client c_0 cannot possibly see player p_1 's blocking move until 125 ms after c_1 initiated it.

This 125 ms minimum reaction window exists regardless of the networking model employed. During this time, if the tick interval is 50 ms, client c_0 will have already made at least two movement decisions without any possibility of knowing about the blocked path. If c_0 was moving toward position (3, 5), they would continue moving toward it for at least two ticks before the blocking becomes visible.

This limitation becomes particularly problematic in competitive scenarios requiring split-second reactions. In a fighting game where a 100 ms reaction time separates amateur from professional players, this network-imposed delay makes true real-time interaction impossible. The only way to reduce this limitation is through better network infrastructure (lower latency), not through improved networking models. This physical constraint

fundamentally shapes how we must design networked games, leading to the prediction and reconciliation techniques discussed in subsequent sections.

3.1.5 Client side prediction

In the (fair) naive model, a client will only be able to apply their input i_t^c to a state $s_{t-t_c^r}$. Therefore a client sees responses to their own input with a latency of t_c^r .

While we previously showed an absolute limitation in seeing other players' inputs, there is no such limitation on seeing state changes to your own inputs. We define a prediction function $f^p : I^c \times S \rightarrow S$ which only progresses the state based on one single client input i_t^c . We write $p_{t,d}^c$ for a client c , predicting from timestamp t , d ticks to the future. Each client now receives state s_t and directly applies their input to get predicted state $p_{t,1}^c = f^p(i_t^c, s_t)$. Since responses to i_t^c will be received in $s_{t+t_c^r}$, we predict the state up to $p_{t,t_c^r}^c$. This method is called client-side prediction and used in most modern games as it provides the client with instantaneous feedback - making gameplay fluent.

The difficulty with this method is to apply new incoming host state s_{t+1} correctly. Given a client predicted $p_{t,t_c^r}^c$, it will receive a response to input i_t^c in at least t_c^r time as the state $s_{t+t_c^r}$. A previous state change arriving as s_{t+1} would override our predicted state $p_{t,t_c^r}^c$, making the game snap, caused by the time difference. The original solution to this problem is to not replace our current state but define a partial application function $f^{\text{partial}} : S \times S \rightarrow S$ which compares the predicted state with the actual state and only replaces diverging parts. During normal operation, the prediction will be close enough causing only rare desynchronizations, which will be visually visible as snapping. We observe, that some parts of the state will always diverge, if state changes are caused by other players or randomness.

One very popular game that is using a slightly modified variation of this method is Minecraft. In Minecraft, the game does not send its direct inputs, but its predicted state. Most important parts of the state to be predicted are breaking blocks and all character movement. The partial application does not happen at client-side, but on host-side. Therefore the host validates if the incoming input looks good enough and doesn't resimulate it if not necessary. When detecting diverging state, the host sends state corrections. Unfortunately this limits anti-cheat capabilities to how sophisticated the validation in the partial application function is.

One major problem with client-side reconciliation is that finding a good partial application function can be difficult. The reason is that we are looking at divergences between s_{t+1}

and p_{t,t_c}^c , therefore a difference in time of $t_c^r - 1 \approx t_c^r$. Therefore even a perfect prediction will diverge on the scale of the round trip time.

One alternative approach would be to not compare the current state for divergence, but memorize past states. Given an interval time of t^i , the client needs to memorize $n_{\text{predict}} = t_c^r/t^i$ game states, which can be partial states. This solution unfortunately doesn't solve the snapping when divergences are found, since it forces the client from a predicted state in time $t + t_c^r$ to $t + 1$.

Example of client-side prediction

To demonstrate client-side prediction in our movement game, we first define a prediction function that allows a client to speculatively apply their own inputs without waiting for server confirmation. Given that we define clients as identifiers, we can formalize the prediction function as follows:

$$\begin{aligned}
 f^{\text{transform}}((i_{\text{up}}, i_{\text{down}}, i_{\text{left}}, i_{\text{right}})) &= ((i_{\text{down}} - i_{\text{up}}, i_{\text{right}} - i_{\text{left}})) \\
 f^{\text{find}}(\text{id}, S) &= x_P, \quad x_P \in \{(\text{id}', x') \in S \mid \text{id} = \text{id}'\} \\
 f^{\text{update}}(\text{id}, s_p, S) &= \{(\text{id}, s_p)\} \cup \{(\text{id}', x') \in S \mid \text{id} \neq \text{id}'\} \\
 f^p(i_c, \text{id}, s) &= f^{\text{update}}(\text{id}, f^{\text{find}}(\text{id}, s) + f^{\text{transform}}(i_c), s)
 \end{aligned}$$

The prediction function f^p takes a client's input i_c , the client's identifier, and the current state, then returns an updated state where only the predicting client's position has changed. Other players remain at their last known positions from the server's authoritative state.

Consider a scenario with client c_0 controlling player p_0 , with a round trip time of 100 ms (50 ms each direction) and a tick interval of 25 ms. Without prediction, the client would experience a 100 ms delay between pressing a key and seeing their character move. With client-side prediction, the response becomes instantaneous:

Tick	Server State	Client Receives	Client Input	Predicted State	Client Display
0	(0, 0), (5, 5)	—	—	—	(0, 0), (5, 5)
1	(0, 0), (5, 5)	—	—	—	(0, 0), (5, 5)
2	(0, 0), (5, 5)	(0, 0), (5, 5)	right	(1, 0), (5, 5)	(1, 0), (5, 5)
3	(0, 0), (5, 5)	(0, 0), (5, 5)	right	(2, 0), (5, 5)	(2, 0), (5, 5)

4	(1, 0), (5, 5)	(0, 0), (5, 5)	right	(3, 0), (5, 5)	(3, 0), (5, 5)
5	(2, 0), (5, 5)	(0, 0), (5, 5)	right	(4, 0), (5, 5)	(4, 0), (5, 5)
6	(3, 0), (6, 5)	(1, 0), (5, 5)	stop	(4, 0), (5, 5)	(4, 0), (5, 5)
7	(3, 0), (7, 5)	(2, 0), (5, 5)	stop	(4, 0), (5, 5)	(4, 0), (5, 5)
8	(3, 0), (8, 5)	(3, 0), (6, 5)	stop	(3, 0), (6, 5)	(3, 0), (6, 5)
9	(3, 0), (9, 5)	(3, 0), (7, 5)	stop	(3, 0), (7, 5)	(3, 0), (7, 5)

At tick 2, the client receives the initial state and immediately begins moving right. Without prediction, the client would display (0, 0) until tick 6. With prediction, the movement appears instantaneous: the client displays (1, 0) at tick 2, even though the server hasn't processed this input yet. The client continues predicting movement, reaching (4, 0) by tick 5.

The critical moment occurs at tick 6 when the client receives the server state (1, 0), (5, 5). This confirms the client's first input was processed, but also reveals that player p_1 has moved to (6, 5) - information the client couldn't predict. At this point, the predicted position (4, 0) must be reconciled with the server's authoritative position (1, 0) plus the three additional predictions made since then.

Notice how the prediction creates a smooth experience for the controlling player: they press right and immediately see movement. However, it also introduces divergence between the predicted and authoritative states. In our example, the client predicted player p_0 at position (4, 0) at tick 6, while the server only had them at (3, 0). This divergence of one position unit represents the unconfirmed prediction still in transit to the server.

The prediction function must carefully handle state updates to avoid "snapping" - sudden position corrections that break immersion. When the server state arrives at tick 8 showing (3, 0), (6, 5), the client has predicted (4, 0) but stopped sending input at tick 6. The reconciliation process must recognize that position (3, 0) represents the final authoritative position after processing all inputs, allowing the display to smoothly converge to the correct state.

This example demonstrates the fundamental trade-off of client-side prediction: immediate responsiveness at the cost of temporary divergence from the authoritative state. The prediction allows fluid gameplay even with 100 ms latency, transforming what would be a sluggish experience into one that feels instantaneous. However, it requires sophisticated reconciliation mechanisms to handle the inevitable corrections when

predicted and authoritative states diverge, particularly when multiple players' actions interact in ways the client cannot predict locally.

3.1.6 Server reconciliation

When merging an update s_{t+1} with a diverged predicted state p_{t,t_c}^c , we previously overrode the state with state s_{t+1} . This skip from $t + t_c^r$ to $t + 1$ causes snapping. What we can do alternatively is to memorize inputs $i_k^c : t + 1 \leq k \leq t + t_c^r$. We replace our state with s_{t+1} and then apply all memorized inputs to get an alternative predicted state $p_{t,t_c}^{c'}$ which would have been reached if the divergence didn't happen. Since we now have a state in time $t + t_c^r$, we do not have snapping behavior. This method is called rollback server reconciliation and is one of the most fundamental techniques in modern games.

Server reconciliation is the process of combining a past received state s_t with a present predicted state p_{t,t_c}^c . Our partial application function is a form of server reconciliation. The goal is to fix divergences in state without the client noticing.

Unfortunately, since we have to rollback to the last received state, we must memorize all predictions we made. In the worst case, we memorize the whole s_t . When receiving s_t , we apply the prediction function n_{predict} times. Therefore an input i_t^c is predicted for the first time at $t - n_{\text{predict}}$ and the last time at t . So each input is processed by the client n_{predict} times. This means that especially clients with longer ping have to do more processing, increasing with a smaller tick interval.

3.1.7 Delta State

The host is sending the whole state s_t each tick. This means we will send redundant data as the state rarely changes completely. We define delta states $\Delta s_t \in \Delta S$ as changes in state which can be applied using a delta application function $f^A : S \rightarrow \Delta S \rightarrow S$. We define an alternative progression function $f^\Delta : I \rightarrow S \rightarrow \Delta S$ returning a delta instead of the whole state and write $f^\Delta(i_t, s_t) = \Delta s_t$. We can derive the original progression function as $f(i_t, s_t) = f^A(s_t, f^\Delta(i_t, s_t))$, therefore all our previous results hold true when using delta states. The other way around we can convert any normal state to delta state given progression function f and state S by defining $f^A = f$, $\Delta S = S$ and $f^A(\cdot, s_t) = s_t$. This means, that in practice we can always implement delta state and use normal state on top.

The host will each tick send Δs_{t+1} to the clients. In the (fair) naive or lockstep model, a client has currently loaded a state s_t and can simply use the delta application function to progress the state. Using rollback reconciliation, the client can use the delta application function after the rollback to s_t and predict the actual state afterwards. If delta state can be used with the partial application reconciliation depends on the state and cannot be generalized because it is not possible to get back to a state s_t .

3.2 Algebraic server reconciliation

3.2.1 Definition

An Abelian group is a set A , together with an operation $+: A \times A \rightarrow A$ which is associative and commutative. Additionally, it has an identity element $a + e = a$ and for every element a an inverse $a + a' = e$.

We assume a delta state that defines an abelian group with $S = \Delta S$, $e = s_e$ and $+= f^A$, where f^A satisfies the required properties. We can observe, that we still have advantages over non-delta state, since we do not need to transfer state, when $\Delta s = s_e$.

We define a delta prediction function $f^{\Delta p}: I^c \times S \rightarrow \Delta S$ producing a delta $\Delta p_{t,d}^c = f^{\Delta p}(i_{t+d}^c, p_{t,d}^c)$ instead of the whole state when predicting the next state. Therefore $p_{t,d+1}^c = p_{t,d}^c + \Delta p_{t,d}^c$ and $p_{t,d}^c = s_t + \sum_{k=0}^d \Delta p_{t,k}^c$.

The client memorizes $n_{\text{predicted}}$ predicted deltas $\Delta p_{t,k}^c: 0 \leq k \leq t_c^r$. We define a new server reconciliation method as follows: For each incoming Δs_{t+1} , we add a correction update $\epsilon = \Delta s_{t+1} - \Delta p_{t,1}^c$ to our current state $p_{t+1,t_c^r}^c = p_{t,t_c^r}^c + \epsilon$. Assuming that predictions are mostly uncorrelated $\Delta p_{t,d}^c \approx \Delta p_{t+1,d-1}^c$, we can show convergence:

$$\begin{aligned}
& p_{t,t_c}^c + \epsilon \\
&= p_{t,t_c}^c + (\Delta s_{t+1} - \Delta p_{t,0}^c) \\
&= s_t + \sum_{k=0}^{t_c} \Delta p_{t,k}^c + (\Delta s_{t+1} - \Delta p_{t,0}^c) \\
&= (s_t + \Delta s_{t+1}) + \sum_{k=1}^{t_c} \Delta p_{t,k}^c + (\Delta p_{t,0}^c - \Delta p_{t,0}^c) \\
&= s_{t+1} + \sum_{k=1}^{t_c} \Delta p_{t,k}^c \\
&\approx s_{t+1} + \sum_{k=0}^{t_c-1} \Delta p_{t+1,k}^c \\
&= p_{t+1,t_c-1}^c
\end{aligned}$$

Therefore we can reach any prediction p_{t+d,t_c}^c over time. We call this reconciliation method algebraic server reconciliation.

3.2.2 Abelian group for ECS games

3.2.3 Limitations

4 Networking Architecture

The goal of a networking system is to provide a seamless interaction with other players inside a game. How to achieve a seamless interaction with other players was investigated in Section 3. The focus of systems architecture is on how to implement these models within the game's underlying system. Consequently, this chapter is of less concern to players and more pertinent to developers. Just as we aim for seamless interaction for players, we also strive for seamless game development for developers. The architecture here is only one possible approach to networking architecture, which will be different depending on needs and game architecture.

The here introduced architecture will give an overview how previously introduced concepts in game networking can be implemented practice and which considerations are made when implementing. It will enable us how our newly introduced method integrates in a networking stack and interacts with other components. It will also give a better understanding in what is happening, when we present our experiments and data afterwards.

4.1 Overview

Our networking system is designed as an independent layer (the “networking layer”) positioned beneath the game system (the “game layer”). The goal is to expose minimal complexity, making game development akin to creating a single-player game. The game system comprises a “world” (representing the game state), game logic that modifies this world, and events. These events facilitate communication between different logic elements and trigger user-facing indicators, such as animations.

This separation of concerns within the game system itself facilitates a cleaner interface with the underlying networking layer, as illustrated below.




Figure 4.1: The fundamental division between the game layer, responsible for gameplay mechanics and presentation, and the network layer, responsible for communication.

4.1.1 Four quadrants model

Visually, the game and networking layers can be divided into four quadrants. The horizontal axis differentiates between the server and the client. As in Section 3, we use the classic server/client model, where the server has authority over the game world. Clients can only interact with the game world by sending their inputs, where the server will send updates to the game world (called replication) and events (for example, to trigger animations). The version of the world maintained on the client, which is partially updated through replication, is referred to as the “world view.”




Figure 4.2: Building on this client-server distinction, we now incorporate the differentiation between the game and network layers.

4.1.2 Server and Client

The vertical axis, as previously introduced, differentiates between the game and network layers. Both the game and networking layers must serve two distinct scenarios: operating as a server or as a client. First, looking at the game layer, the server usually does not want to render the world. This independence not only boosts performance by avoiding unnecessary rendering overhead on the server but also simplifies porting the server logic to environments without graphical capabilities.

Most game logic is also only needed on the server, as changes are replicated authoritatively to the client. An exception could be user-controlled elements of the game world, such as the player character. However, similar to the approach in Section 3, we distinctly define this as “prediction logic” to improve cohesion. Additional event-based visuals, like animations, are typically only required on the client, as they do not impact the core game logic. This ensures that client-side presentation details do not impose constraints or unnecessary complexity on the authoritative server-side game simulation.




Figure 4.3: The server manages the authoritative World and Logic, while the client maintains a World View, handles Prediction Logic, Rendering, and event-based Animations.

4.1.3 Game and Network

The networking layer must cater to the server logic by providing a list of player inputs each game tick. Before replication, the networking layer has to look for changes in the server world. As introduced with “delta state” in Section 3, the system aims to send only the changes in the game world (deltas) rather than the entire world state each time. The efficiency of this delta state transmission is crucial for minimizing bandwidth usage and server load.




Figure 4.4: The complete four-quadrant architecture, illustrating the interplay between Game and Network layers on both Server and Client, including data flows.

Depending on how the world in the server is structured, it might provide additional helpers to detect changes. When the client receives changes, it applies them to its local game world, provided these changes are not already predicted. Otherwise, we need to use reconciliation to update our world view. This continuous loop of input collection by the network layer, server processing, state replication, and client-side reconciliation (if necessary) forms the core of the real-time networked experience.

4.2 Server networking layer

Previously, we explained an overview of our architecture and how components abstractly interact. Here, we want to dive deeper into the server-side dataflow. The following diagram displays all relevant components of the server-side networking layer, including their interaction with the game state.




Figure 4.5: Server-side networking components and dataflow, illustrating the processing of incoming inputs and the generation of outgoing changes and events.

We have three primary communication channels with clients: incoming inputs from players, outgoing changes to replicate the game world, and outgoing events. Events are not strictly part of the game world but are needed for richer visual experiences on the client. The input tracker collects incoming inputs from players, providing a consolidated input to the game logic each tick. It is also responsible for providing stale or empty inputs if a new input has not arrived in time. The change detection mechanism is responsible for replicating the game world to all clients, ensuring each client has the same (or a consistent partial) world view by exchanging deltas.

The dataflow can be summarized by the following steps, which correspond to the numbered arrows in the diagram:

1. Users send their inputs to the server. The server receives these inputs, and the input tracker keeps track of them. Since inputs can be sent via unreliable communication methods, this component might also include logic for filtering inputs (e.g., processing only the most recent ones) or handling out-of-order packets.
2. Each tick, a set of all relevant inputs (one per player, potentially synthesized by the input tracker if an actual input is missing) is provided to the game logic. The game logic itself is fully defined in the game layer but is driven by inputs supplied by the networking layer.
3. Using these inputs, the game logic in the game layer modifies the world state. Optionally, the game logic may directly inform the change detection component about specific changes that have occurred.
4. Modifications to the world state need to be communicated to all clients. This is managed by the change detection component, which generates deltas representing these modifications. Since world state changes happen frequently and quickly, unreliable communication methods might be preferred for sending these deltas to clients. However, using unreliable transport for state updates adds complexity, such as the need for the server (or client) to track applied changes and potentially re-send missed updates to ensure eventual consistency.
5. Processing the game logic can also trigger occurrences that are not direct modifications of the persistent game world but are important for the player experience (e.g., sound effects, temporary visual effects). These are transmitted directly to the corresponding clients as events. Events are often sent reliably if they are critical, or unreliably if they are transient and missing one isn't detrimental.

This entire process is repeated each game tick while the game is running.

4.3 Client networking layer

The overall structure and communication pathways of the client-side networking layer were introduced in the architectural overview. Having inspected the dataflow within the server-side networking layer, we now turn our attention to the client-side components.

This part is particularly important for our discussion, as it is where our new algebraic reconciliation method integrates. The following diagram presents the client-side dataflow.




Figure 4.6: Client-side networking components and dataflow, detailing the processing of local player inputs, incoming server changes and events, and the interaction between prediction and reconciliation modules.

Similar to the server-side, the client has three primary communication channels: incoming events from the server, incoming state changes from the server, and outgoing player input. We will not focus extensively on event handling in this section, assuming that the world view component processes them (e.g., by triggering local visual or audio effects). The reconciliation, prediction, and prediction logic components collectively form the most interesting part of this architecture for our purposes. The interactions between these three

elements differ depending on the specific reconciliation strategy employed, as will be detailed in the following subsections. We assume the client's representation of the game world is conceptually split into a "normal" part (the world view, reflecting acknowledged server state) and a predicted part (reflecting speculative local changes).

The dataflow can be summarized by the following steps, which correspond to the numbered arrows in the diagram:

1. Each tick, the server replicates state changes (deltas) to the client. These changes can correspond to either the predicted or non-predicted parts of the client's world representation. If a change affects a non-predicted part, it is directly applied to the world view. Otherwise, if it affects a predicted part, it is passed to the reconciliation module.
2. In parallel with receiving world state changes, the client also receives events from the server. As mentioned, we assume the world view handles these, for instance, by initiating animations or sound effects that do not alter the core game state logic.
3. Each tick, the game layer captures inputs from the player. These inputs are sent to the server as quickly as possible, often using an unreliable communication channel to minimize latency. Critically for client-side responsiveness, these inputs are also passed to the local prediction module.
4. After the prediction module receives inputs from the game layer, these inputs are fed into the game-specific prediction logic. If rollback reconciliation is being used (as detailed in Section 4.3.1), this prediction module is also responsible for memorizing these inputs for potential future re-simulation.
5. The game-layer defined prediction logic uses the local inputs to modify the predicted part of the world, most commonly affecting entities directly controlled by the player, such as their character. This immediate local modification provides the player with responsive feedback to their actions.
6. When incoming state changes from the server (received in step 1) affect parts of the world that have been locally predicted, these discrepancies must be handled by the reconciliation module. This module receives the authoritative server change and information about the current predicted state from the prediction module.
7. If rollback reconciliation is employed, the reconciliation module will use the memorized inputs (from step 4) and the authoritative server state to re-simulate the player's actions and compute a corrected predicted part of the world. For other reconciliation methods, like algebraic reconciliation (detailed in Section 4.3.2), this



step will differ. After processing, the reconciliation module updates the predicted part of the world to align it more closely with the authoritative server state, aiming to correct any mispredictions smoothly.

Similarly to the server-side dataflow, this entire process is repeated each game tick while the game is running.

4.3.1 Rollback reconciliation

When an incoming authoritative state change (delta) from the server affects the predicted part of the client's game world, reconciliation is necessary to correct any mispredictions. The most common reconciliation method currently used in games is rollback reconciliation, which was introduced conceptually in Section 3.1.6. The core idea behind rollback reconciliation is to revert the client's game state to the last known authoritative state before the misprediction occurred, apply the incoming server correction, and then re-simulate all local player inputs that have occurred since that authoritative state.

The interactions for this process within our client-side architecture are illustrated below.




Figure 4.7: Dataflow for rollback reconciliation on the client. Incoming changes trigger a rollback, replay of memorized inputs via the prediction module, and update of the predicted world.

The primary challenge in implementing rollback reconciliation lies in efficiently reverting the game state. When working with delta states (as discussed in Section 3.1.7), two common approaches are:

-
1. **State Cloning:** Maintain at least two distinct copies (or snapshots) of the predicted part of the world. One copy represents the state before the latest batch of local predictions is applied. When an incoming server change necessitates a rollback, the current predicted world (with the mispredictions) is discarded, and the client reverts to the previous clean snapshot. The server change is applied to this snapshot, which then becomes the new baseline for re-applying subsequent local inputs.
 2. **Storing Reverse Deltas:** For each local prediction applied to the predicted part of the world, store a corresponding “reverse delta” or “undo operation.” To roll back, these reverse deltas are applied in reverse chronological order to the current predicted state, effectively reverting it step-by-step to the desired past authoritative state. Once rolled back, the server’s delta is applied, and then local inputs are re-simulated.

After successfully rolling back the state to the point of divergence, the reconciliation module needs to re-apply the local inputs that the player has generated since that point. To facilitate this, the prediction module, as mentioned in the client-side dataflow (see step 4), memorizes these applied inputs in a buffer. The rollback reconciliation process accesses this buffer to retrieve the sequence of inputs that must be re-played by the prediction logic against the corrected state.

4.3.2 Algebraic reconciliation

A significant drawback of rollback reconciliation, as highlighted in Section 3.1.6, is the computational cost of re-simulating predictions. Each time an authoritative server update corrects the client’s predicted state, the client might need to re-process multiple ticks of input. This re-simulation effort can increase with the client’s network latency, potentially affecting performance, especially in games with many frequently predicted game objects or for players with higher ping times.

The dataflow for applying algebraic reconciliation within our client architecture is depicted in the figure below.




Figure 4.8: Dataflow for algebraic reconciliation on the client. Incoming server deltas are used by the reconciliation module to calculate a direct correction delta.

Our proposed algebraic server reconciliation method, whose theoretical basis was introduced in Section 3.2, offers an alternative designed to avoid this repetitive prediction overhead. This method relies on representing game state changes as “delta states.” For algebraic reconciliation to work effectively, these delta states need to possess certain well-defined mathematical properties, specifically those of an Abelian group, as detailed in our theory section. In essence, this means the deltas can be reliably combined, effectively undone, and the sequence of their combination does not alter the final result, much like

arithmetic operations on numbers. When these conditions are met, corrections to the client's predicted state can often be calculated and applied as a single, direct adjustment, rather than requiring a full re-simulation of past inputs.

The implementation of algebraic reconciliation within the client-side networking layer primarily modifies how the reconciliation module processes incoming server changes and updates the client's predicted part of the world. The client's prediction logic, when handling local player inputs, conceptualizes its effects as a series of "predicted deltas." The prediction module tracks the cumulative outcome of these locally generated deltas, effectively maintaining the client's speculative view of the game state built upon the last acknowledged server state.

Concurrently, the client receives authoritative state updates from the server, also in the form of deltas. When such an authoritative server delta arrives, the reconciliation module's task is to compute a "correction delta." This is achieved by comparing the server's authoritative delta for a given time period with the net predicted delta that the client had generated for that same period. The discrepancy between these two deltas isolates the misprediction. This "difference," which is itself a delta, forms the required correction. The ability to reliably calculate this correction delta hinges on the algebraic (Abelian group) properties of the delta states, which allow for a meaningful "subtraction" or "inversion" of the predicted deltas relative to the server's authoritative ones, as outlined in Section 3.2.

Once this single correction delta is determined, it is applied directly to the client's current predicted part of the world. This direct application adjusts the client's state to more accurately reflect the server's view, effectively integrating the server's information without the need to rewind and individually re-process past local inputs. This avoidance of extensive input re-simulation is a key advantage over the rollback method. While the re-simulation of old inputs is bypassed, the client naturally continues to apply new local inputs via the prediction logic to its newly corrected state to maintain immediate responsiveness. It's important to note that for this process to function, the prediction module must maintain a record of the predicted deltas it has generated, as this information is essential for calculating the correction delta. This is distinct from rollback reconciliation's requirement to store the raw inputs themselves for potential re-simulation.

By directly adjusting the predicted state through these well-behaved delta manipulations, the algebraic approach can offer a more computationally efficient reconciliation process. This is particularly beneficial when the game state can be effectively represented by deltas that adhere to the required algebraic properties, as this allows for the direct calculation and application of corrections.



5 Experiments



6 Discussion



7 Future work



Bibliography

- [1] “Das Bild der TU Darmstadt.” Accessed: May 1, 2020. [Online]. Available at: https://www.intern.tu-darmstadt.de/media/medien_stabsstelle_km/services/medien_cd/das_bild_der_tu_darmstadt.pdf